

Modeling Wine Quality Using Physicochemical Attributes

Predictive Classification, Enological Feature Analysis, and Operational Decision Support

Author: Anthony S. Candelas (Lead Modeling Engineer / Author)

Academic Affiliation: University of San Diego, Shiley-Marcos School of Engineering

Program: Master of Science in Applied Artificial Intelligence (MS-AAI)

Date: October 2026

Repository: <https://github.com/AnthonyCandelas/MAAI-USD-ASC>

Executive Abstract

In commercial viticulture and global wine distribution, wine quality has historically been assessed through post-bottling sensory evaluation by professional tasting panels. While human expertise provides rich qualitative descriptions, sensory evaluation is inherently subjective, logistically constrained, and unscalable across commercial production volumes. This investigation evaluates whether laboratory-measured physicochemical attributes can predict expert quality ratings of Portuguese Vinho Verde wines.

Using a cleaned cohort of 5,320 observations (1,359 red and 3,961 white wines) and 12 feature variables (11 laboratory chemical measurements plus wine type), we frame quality prediction as a cost-sensitive classification problem separating high-quality wines (expert sensory score 7 or higher, 18.98% prevalence) from ordinary wines (score 6 or lower, 81.02% prevalence).

Using a 10-fold stratified cross-validation architecture on the training set ($N = 4,256$), a bagging Random Forest classifier with balanced class weighting demonstrated statistically significant superiority over a regularized Logistic Regression baseline (paired t-test: $t = 3.3840$, $p = 0.0081$; Nadeau & Bengio corrected $t = 2.3290$, $p = 0.0448$) and outperformed gradient boosting. When evaluated on the completely unseen holdout test dataset ($N = 1,064$), the final Random Forest model achieved an **Accuracy of 83.08%**, a **Sensitivity (Recall) of 71.78%**, an **F1-Score of 0.6170**, an **ROC-AUC of 0.8738**, and a **PR-AUC of 0.6230** (substantially exceeding the 0.1898 no-skill baseline). Subgroup disaggregation confirmed strong discriminability across varieties, achieving **90.18% accuracy and 0.9307 ROC-AUC on red wines**.

Feature importance evaluations (Gini impurity and permutation F1 drop) revealed that **Alcohol content (+0.2461 test F1 drop)** is the dominant physiological marker of grape maturity and extraction, while **Chlorides (+0.0796 F1 drop)** and **Density (+0.0768 F1 drop)** serve as primary structural boundaries against bitterness and imbalance. We outline a three-tier operational deployment model for wineries to automate early batch screening, adjust operating thresholds based on brand risk, and prioritize expert sommelier panel evaluations where they provide the greatest economic value.

1. Project Background and Objective

The commercial success of winemaking depends heavily on product consistency and accurate quality tier classification. Premium wines command substantial retail price margins, whereas ordinary table wines are marketed for volume consumption. However, chemical measurements alone are rarely integrated directly into commercial decision-making because individual chemical variables exhibit non-linear interactions with human sensory perception.

The primary objective of this investigation is to construct, validate, and analyze a machine learning classification framework that predicts whether a wine will achieve a high-quality expert rating (quality score 7 or higher) based solely on laboratory chemistry. The analysis addresses three core questions:

1. Which machine learning architecture provides the optimal balance of sensitivity, precision, and generalizability on imbalanced enological data?
2. Which specific physicochemical attributes most strongly govern wine quality?
3. How can wineries and distributors deploy calibrated probabilistic outputs to optimize blending, batch screening, and premium bottling decisions?

2. Data Cleaning and Preparation

Data Sources

The data comes from the Wine Quality dataset in the UCI Machine Learning Repository (Cortez et al., 2009). It has two files, one for red wines (1,599 samples) and one for white wines (4,898 samples), all from the Vinho Verde region of Portugal. Each sample has 11 physicochemical measurements taken in the lab: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, and alcohol. Each wine also has a quality score from 0 to 10, which is the median of at least three ratings from wine experts. In practice, the scores only range from 3 to 9.

Both files are semicolon-separated, so we set the separator when loading them with pandas.

Combining Red and White Wines

We combined the red and white files into a single dataset of 6,497 rows instead of analyzing them separately. This gives the models more data to learn from, and it lets us compare the two wine types directly. To keep track of where each row came from, we added a `type` column ("red" or "white") before combining. We also renamed the columns to `snake_case` (for example, `fixed acidity` became `fixed_acidity`) so they are easier to use in code.

Data Types and Missing Values

All 11 measurements loaded as floating-point numbers, `quality` loaded as an integer, and `type` as text. These are the types we expected, so no conversions were needed.

We found no missing values in any column, so no rows had to be removed or filled in for this reason.

Duplicate Records

We found 1,177 rows that were exact copies of another row, meaning every measurement, the quality score, and the wine type all matched. This was 240 red wines and 937 white wines, about 18% of the combined data.

We decided to remove the duplicates and keep only the first copy of each. With 11 measurements recorded to several decimal places, it is unlikely that this many different wines would match on every value, so most of these rows are probably repeated entries of the same sample. Keeping them would also cause a problem later: when the data is split, identical rows could end up in both the training and test sets. The model would then be tested partly on data it had already seen, and its performance would look better than it really is. After removing duplicates, 5,320 wines remained (1,359 red and 3,961 white).

Outliers

We checked each measurement for outliers using the interquartile range (IQR) rule, where a value is an outlier if it falls more than 1.5 times the IQR below the first quartile or above the third quartile. Table 1 shows the results.

Table 1. Outliers per feature using the $1.5 \times \text{IQR}$ rule (after removing duplicates, $n = 5,320$)

Feature	Outliers	Feature	Outliers
Fixed acidity	304	pH	49
Volatile acidity	279	Free sulfur dioxide	44
Chlorides	237	Total sulfur dioxide	10
Sulphates	163	Density	3
Citric acid	143	Alcohol	1
Residual sugar	141		

In total, 1,094 wines had at least one outlying value. We also compared each feature across red and white wines using box plots (see Appendix).

We decided to keep all of the outliers, for three reasons:

1. **Many of them come from combining two different types of wine.** Red and white wines have different chemistry. For example, the median volatile acidity is 0.52 g/L for red wines and 0.26 g/L for white wines, and the median total sulfur dioxide is 38 mg/L for red and 133 mg/L for white. A value that looks extreme for the combined data can be normal for its own wine type.
2. **The extreme values are realistic.** The highest residual sugar (65.8 g/L) matches a sweet wine, and the highest sulfur dioxide values (289 mg/L free, 440 mg/L total) are high but possible. None of the values are impossible, such as negative concentrations, so they are most likely real wines and not data entry errors.
3. **Removing them would lose information.** About one in five wines has at least one outlier, so dropping them would shrink the dataset a lot and could remove exactly the unusual wines that are useful to learn from. If outliers become a problem for a specific model, feature scaling or tree-based models can handle them.

Target Variable

Since our goal is classification, we turned the quality score into a binary target called `is_good`. A wine is labeled good (`is_good = 1`) if its quality is 7 or higher, and not good (`is_good = 0`) otherwise. After cleaning, 1,009 wines (19.0%) are labeled good and 4,311 (81.0%) are not good. Good wines are more common among white wines (20.8%) than red wines (13.5%).

The classes are imbalanced. A model that always predicts "not good" would already be about 81% accurate, so accuracy alone is not a good way to judge our models. We use precision, recall, F1 score, and ROC-AUC in the model evaluation instead.

We also added a numeric column `is_red` (1 for red, 0 for white) so the models can use wine type as an input. The original `quality` column is not used as a model input, because `is_good` is calculated directly from it.

Train/Test Split

We split the cleaned data into a training set (80%, 4,256 wines) and a test set (20%, 1,064 wines). The split was stratified on `is_good`, so both sets have the same share of good wines (19.0%). We used a fixed random seed (42) so every team member gets the same split and the results can be reproduced.

We did not scale the features at this stage. A scaler has to be fit only on the training data, otherwise information from the test set leaks into training. Scaling is done as part of the modeling step instead.

Summary

Table 2. Summary of cleaning steps

Step	Result
Load and combine	6,497 wines (1,599 red, 4,898 white), <code>type</code> column added
Data types	All as expected, no changes
Missing values	None found
Duplicates	1,177 removed, 5,320 wines remain
Outliers	Found with the IQR rule and kept
Target	<code>is_good</code> = 1 if quality $\geq$ 7 (19.0% of wines)
Train/test split	80/20, stratified on <code>is_good</code> , 4,256 train and 1,064 test

Reference

Cortez, P., Cerdeira, A., Almeida, F., Matos, T., & Reis, J. (2009). Modeling wine preferences by data mining from physicochemical properties. *Decision Support Systems*, 47(4), 547-553.

3. Exploratory Data Analysis

We explored the full cleaned dataset (`wine_clean.csv`, 5,320 wines) to understand how each feature is distributed, how red and white wines differ, and which features are related to wine quality. The goal was to find out which features are likely to be useful for modeling and which problems the models would need to handle.

Summary Statistics


Table 3 shows summary statistics for the 11 chemical features. The features are measured on very different scales. Density varies by less than 0.06 g/cm³ across all wines, while total sulfur dioxide ranges from 6 to 440 mg/L. For several features the mean is noticeably larger than the median, especially residual sugar (5.05 vs. 2.70 g/L), which is an early sign of right skew.

Table 3. Summary statistics of the chemical features (n = 5,320)

Feature	Mean	Std. dev.	Min	Median	Max
Fixed acidity (g/L)	7.215	1.320	3.800	7.000	15.900
Volatile acidity (g/L)	0.344	0.168	0.080	0.300	1.580
Citric acid (g/L)	0.318	0.147	0.000	0.310	1.660
Residual sugar (g/L)	5.048	4.500	0.600	2.700	65.800
Chlorides (g/L)	0.057	0.037	0.009	0.047	0.611
Free sulfur dioxide (mg/L)	30.037	17.805	1.000	28.000	289.000
Total sulfur dioxide (mg/L)	114.109	56.774	6.000	116.000	440.000
Density (g/cm ³)	0.995	0.003	0.987	0.995	1.039
pH	3.225	0.160	2.720	3.210	4.010
Sulphates (g/L)	0.533	0.150	0.220	0.510	2.000
Alcohol (% vol.)	10.549	1.186	8.000	10.400	14.900

Distribution of the Target

Figure 1 shows the original quality scores. Most wines received a 5 (1,752 wines) or a 6 (2,323 wines), and very few received the extreme scores of 3 (30 wines) or 9 (5 wines). Using quality ≥ 7 as the cutoff, 1,009 wines (19.0%) are good and 4,311 (81.0%) are not good. Good wines are more common among white wines (20.8%) than red wines (13.5%).



f01_quality_scores.png


Figure 1. Distribution of quality scores. Scores of 7 and above are labeled good.

The imbalance means a model that always predicts "not good" would be 81% accurate while being useless. This confirms that accuracy should not be the main metric for comparing models.

Feature Distributions

Figure 2 shows a histogram of each feature, and the skewness values confirm what the plots show. Most features are right-skewed, with a long tail of high values. Chlorides is by far the most skewed (skewness 5.34), followed by sulphates (1.81), residual sugar (1.71), fixed acidity (1.65), volatile acidity (1.50), and free sulfur dioxide (1.36). pH, citric acid, alcohol, and density are only mildly skewed.

Total sulfur dioxide has a skewness close to zero (0.06), but its histogram has two peaks. This is not a single symmetric distribution. It is two distributions on top of each other, one for red wines and one for white wines, as the next section shows.




f02_feature_histograms.png

Figure 2. Histograms of the 11 chemical features.

Skewed features with very different scales can cause problems for models that rely on distances or linear combinations of the inputs, such as logistic regression. For these models we standardize the features. Tree-based models split on one feature at a time and are not affected by skew or scale.

Red vs. White Wines

Red and white wines have very different chemistry (Figure 3). The biggest differences are in sulfur dioxide and sugar. White wines have a median total sulfur dioxide of 133 mg/L compared with 38 mg/L for red wines, more than twice as much free sulfur dioxide (33 vs. 14 mg/L), and more than twice as much residual sugar (4.7 vs. 2.2 g/L). Red wines have about twice the volatile acidity (0.52 vs. 0.26 g/L) and chlorides (0.079 vs. 0.042 g/L), and higher sulphates and fixed acidity. Alcohol, density, and pH are similar for both types.



f03_red_vs_white.png


Figure 3. Each feature by wine type.

These differences explain many of the outliers found during data cleaning. A value that is normal for one type of wine can look extreme in the combined data. Their quality scores, on the other hand, are similar: the average quality is 5.85 for white wines and 5.62 for red wines. Because wine type affects so many features, we kept `is_red` as a model input so the models can take these differences into account.

Correlations

Figure 4 shows the correlation matrix of all numeric variables. Two pairs of features are strongly correlated with each other:

- **Free and total sulfur dioxide** ($r = 0.72$). Free sulfur dioxide is part of the total, so this is expected.
- **Density and alcohol** ($r = -0.67$). Alcohol is less dense than water, so wines with more alcohol are lighter. Density is also related to residual sugar ($r = 0.52$), since dissolved sugar makes wine denser.




f04_correlation_heatmap.png

Figure 4. Correlation matrix of all numeric variables.

This multicollinearity does not hurt a model's predictions much, but it can make the coefficients of a linear model unstable and hard to interpret, because the model cannot tell which of two related features is responsible for an effect. We keep this in mind when interpreting the logistic regression model in the Model Analysis section.

Figure 5 shows how each feature is correlated with our target, `is_good`. (The `quality` column is left out because `is_good` is calculated from it.) **Alcohol has by far the strongest relationship with being a good wine ($r = 0.42$).** Density ($r = -0.30$), chlorides ($r = -0.16$), and volatile acidity ($r = -0.14$) have the strongest negative relationships. The remaining features all have correlations close to zero ($|r| < 0.09$). On their own they say little about whether a wine is good, although they may still help in combination with other features.



f05_target_correlation.png

Figure 5. Correlation of each feature with `is_good`.


Features by Quality Class

Table 4 compares the average feature values of good and not good wines. Good wines have more alcohol (11.57% vs. 10.31%), lower volatile acidity (0.294 vs. 0.356 g/L), lower chlorides (0.044 vs. 0.060 g/L), lower density, and less residual sugar and total sulfur dioxide. Sulphates, pH, citric acid, and free sulfur dioxide are almost the same in both groups.

Table 4. Average feature values by quality class

Feature	Not good	Good
Fixed acidity (g/L)	7.247	7.079
Volatile acidity (g/L)	0.356	0.294
Citric acid (g/L)	0.314	0.337
Residual sugar (g/L)	5.233	4.259
Chlorides (g/L)	0.060	0.044
Free sulfur dioxide (mg/L)	29.939	30.453
Total sulfur dioxide (mg/L)	115.958	106.211
Density (g/cm ³)	0.995	0.993
pH	3.221	3.241
Sulphates (g/L)	0.531	0.545
Alcohol (% vol.)	10.310	11.573

Figure 6 shows the clearest difference, alcohol. The middle half of good wines falls between 10.8% and 12.4% alcohol (median 11.6%), while the middle half of not good wines falls between 9.5% and 11.0% (median 10.1%). There is still overlap between the groups, so alcohol alone cannot separate good wines from the rest. Box plots of volatile acidity and sulphates, and scatterplots of alcohol and volatile acidity against the quality score, are in the Appendix.



f06_alcohol_boxplot.png

Figure 6. Alcohol content of good and not good wines.

All of these relationships are associations. They do not show that changing a feature would change a wine's quality.

Key Takeaways for Modeling

1. **Class imbalance.** Only 19% of wines are good. Models should be trained with class weights and judged on precision, recall, F1, ROC-AUC, and PR-AUC rather than accuracy.
2. **Most useful features.** Alcohol is the strongest single predictor, followed by density, chlorides, and volatile acidity. No single feature separates the classes well, so a model that combines features is needed.
3. **Skew and scale.** Most features are right-skewed and on very different scales, so logistic regression needs standardized inputs. Tree-based models do not.
4. **Multicollinearity.** Free and total sulfur dioxide, and density, alcohol, and residual sugar, are correlated. This matters most when interpreting linear model coefficients.
5. **Wine type.** Red and white wines have very different chemistry, so `is_red` is kept as an input.

4. Model Selection and Architecture Benchmarking

In this phase, we built and evaluated machine learning classifiers to predict whether a wine is high quality based on its 11 physicochemical measurements and wine type. Following the definition from data cleaning, a wine is good (`is_good = 1`) if its sensory evaluation score is 7 or higher, and ordinary (`is_good = 0`) if its score is 6 or lower.

The feature set includes 12 variables:

- 11 continuous chemical features: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, and alcohol.
- 1 binary indicator: `is_red` (1 for red wine, 0 for white wine).

To prevent target leakage, the original continuous `quality` score and the text `type` column were dropped before modeling.

Preprocessing and Prevention of Data Leakage

Data leakage happens when information from outside a training fold is used to transform features or select parameters. To prevent leakage:

1. **Strict Partitioning:** The 80/20 stratified split from data cleaning was strictly preserved. All model exploration, scaling, parameter estimation, and cross-validation were conducted exclusively on `train.csv` (N = 4,256). The test partition (`test.csv` , N = 1,064) remained completely untouched until the final evaluation.
2. **Encapsulated Preprocessing Pipelines:** For distance-sensitive algorithms requiring feature normalization (specifically Logistic Regression), standard z-score normalization (`StandardScaler`) was encapsulated within a scikit-learn `Pipeline` . During 10-fold cross-validation, the scaler was fit solely on the 9 training folds in each split and then applied to transform the held-out validation fold. Tree-based models (Random Forest and Gradient Boosting) do not require feature scaling, so they operate directly on the raw chemical values.

Handling Class Imbalance

As documented in data cleaning, the target distribution is notably imbalanced:

- Ordinary Wines (`is_good = 0`): 81.04% of training observations (3,449 samples).
- High-Quality Wines (`is_good = 1`): 18.96% of training observations (807 samples).

A baseline classifier that always predicts the majority class would achieve 81.04% accuracy while failing completely to detect a single high-quality wine. To address this without distorting the data with synthetic oversampling:

1. **Stratified Partitioning:** A 10-fold Stratified K-Fold cross-validation scheme (`StratifiedKFold(n_splits=10, shuffle=True, random_state=42)`) was used for all model comparisons. This guarantees that each validation fold contains the exact 18.96% positive class proportion.
2. **Cost-Sensitive Weighting:** All candidate models incorporated cost-sensitive loss weighting (`class_weight='balanced'`). The loss associated with misclassifying a minority-class observation was scaled inversely to class frequency according to:

$$w_j = \frac{1}{2 \times N_j}$$

where N is the total number of training samples and N_j is the count of samples in class j . This assigns a weight of approximately 2.64 to high-quality wines and 0.62 to ordinary wines, penalizing false negatives and compelling the algorithms to balance precision and recall.

Evaluation Metrics

Because the classes are imbalanced, classification accuracy is uninformative. We evaluate each model using five metrics:

1. **Precision:** $\frac{TP}{TP + FP}$. The share of predicted high-quality wines that are truly high quality. High precision protects brand reputation by preventing ordinary wines from being bottled as premium reserves.
2. **Recall (Sensitivity):** $\frac{TP}{TP + FN}$. The share of actual high-quality wines successfully identified by the model. High recall minimizes lost revenue by preventing premium wines from being sold as table wine.
3. **F1-Score:** The harmonic mean of precision and recall ($2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$). F1-score serves as our primary selection metric because it balances precision and recall without sacrificing either one.
4. **ROC-AUC:** Area Under the Receiver Operating Characteristic curve. Measures overall ranking and discrimination across all classification thresholds, independent of class prevalence.
5. **PR-AUC (Average Precision):** Area Under the Precision-Recall curve. Focuses specifically on the minority class trade-off across all thresholds, evaluated against the 0.1896 random baseline.

Candidate Models

We benchmarked three model families to find the best balance of predictive power, resistance to overfitting, and interpretability:

1. **Logistic Regression (Parametric Baseline):**
 - *Rationale:* A linear reference baseline. It models log-odds as a linear combination of normalized inputs with L_2 ridge regularization ($C = 1.0$) and balanced class weights.
 - *Limitations:* Assumes linear decision boundaries and cannot capture non-linear chemical interactions (such as synergies between alcohol, sugar, and volatile acidity) without manual interaction terms.

2. Random Forest Classifier (Ensemble Bagging: Primary Model):

- *Rationale:* An ensemble of 300 decorrelated decision trees built on bootstrap samples with random feature subsampling. Random forest handles skewed chemical distributions, resists outliers, captures non-linear interactions, and mitigates single-tree variance through bagging.
- *Regularization:* Constrained to a maximum tree depth of 15, minimum split size of 4, and minimum leaf size of 2 to prevent memorizing noise in expert ratings.

3. Histogram-Based Gradient Boosting (Ensemble Boosting: Comparative Model):

- *Rationale:* An iterative sequential boosting model that bins continuous features into 256 integer histograms, optimizing logistic loss through gradient descent.
- *Configuration:* Evaluated with 200 boosting iterations, a learning rate of 0.08, maximum tree depth of 6, and balanced sample weighting.

Cross-Validation Benchmark


Table 5 summarizes the cross-validation performance of all three candidate architectures across the 10 stratified validation folds. All metrics report the fold mean plus or minus one standard deviation.

Table 5. 10-Fold Stratified Cross-Validation Benchmark on Training Dataset (N = 4,256)

Model Architecture	F1-Score	ROC-AUC	PR-AUC	Recall	Precision	Accuracy
Logistic Regression (Baseline)	0.5316 +/- 0.0263	0.8306 +/- 0.0195	0.5282 +/- 0.0461	0.7869 +/- 0.0250	0.4021 +/- 0.0282	0.7361 +/- 0.0251
Random Forest (Final Model)	0.5599 +/- 0.0338	0.8492 +/- 0.0155	0.5747 +/- 0.0423	0.6543 +/- 0.0501	0.4908 +/- 0.0352	0.8050 +/- 0.0170
HistGradientBoosting	0.5473 +/- 0.0216	0.8386 +/- 0.0181	0.5574 +/- 0.0342	0.6419 +/- 0.0283	0.4779 +/- 0.0273	0.7984 +/- 0.0140

Note: All candidate models utilize balanced class weighting and reproducible seed (random_state = 42). Bold values indicate the superior mean score in each metric column.

Figure 7 illustrates the comparative performance across all six validation metrics.



figure_01_model_cv_comparison.png

Key Performance Insights:

- **Baseline Trade-off:** Logistic Regression achieved high recall (78.69%) but low precision (40.21%), generating excessive false positive classifications. Its overall F1-score (0.5316) was the lowest among candidate models.
- **Ensemble Superiority:** Random Forest achieved the highest F1-score (0.5599), ROC-AUC (0.8492), PR-AUC (0.5747), and overall accuracy (80.50%). Its precision (49.08%) exceeded the baseline by nearly 9 percentage points.
- **Boosting Comparison:** HistGradientBoosting performed competitively (F1 = 0.5473, ROC-AUC = 0.8386) but did not surpass Random Forest in any key metric.

Statistical Comparison

To test whether Random Forest's higher F1-score was statistically significant or merely due to random split variation, we evaluated the differences across the 10 paired validation folds (Table 6).

Table 6. Statistical Significance Testing on 10-Fold Cross-Validation F1-Scores

Comparison	Mean F1 Diff	Standard Paired t (p-value)	Nadeau & Bengio Corrected t (p-value)	Statistical Decision (alpha = 0.05)
Random Forest vs. Logistic Regression	+0.0283	t = 3.3840 (p = 0.0081)	t = 2.3290 (p = 0.0448)	Statistically Significant (p < 0.05)
Random Forest vs. HistGradientBoosting	+0.0126	t = 1.4398 (p = 0.1838)	t = 0.9908 (p = 0.3477)	Not Significant at alpha = 0.05

Statistical Findings:

1. **Comparison with Linear Baseline:** The standard paired t-test between Random Forest and Logistic Regression yields $t = 3.3840$ ($p = 0.0081$). Because standard paired tests on cross-validation folds can overstate significance due to overlapping training folds, we also calculated the Nadeau and Bengio (2003) corrected

resampled t-test, which incorporates the variance correction factor $(1/k + n_{\text{test}}/n_{\text{train}}) = 1/10 + 1/9 \approx 0.2111$). Under this conservative test, the improvement remains statistically significant ($t = 2.3290$, $p = 0.0448 < 0.05$). We reject the null hypothesis of equal performance at the 95% confidence level, confirming that Random Forest provides a genuine improvement over the linear baseline.

2. **Comparison with Boosting:** The paired t-test between Random Forest and HistGradientBoosting yields $t = 1.4398$ ($p = 0.1838$; Nadeau & Bengio corrected $t = 0.9908$, $p = 0.3477$). While this difference is not statistically significant at $\alpha = 0.05$, Random Forest consistently achieved a higher mean F1-score (+0.0126), higher ROC-AUC (+0.0106), higher PR-AUC (+0.0173), and better precision (+0.0129).

Hyperparameter Tuning

We ran a grid search over key hyperparameters on the training set using 10-fold cross-validation. Table 7 presents the top five parameter configurations ranked by validation F1-score.

Table 7. Hyperparameter Sensitivity Analysis for Random Forest (10-Fold CV F1-Score)

Rank	Max Depth	Min Leaf Size	Trees (Estimators)	Validation F1 (Mean +/- Std)	Training F1 (Mean)
1	15	2	300	0.5599 +/- 0.0338	0.9372
2	20	2	300	0.5584 +/- 0.0346	0.9615
3	15	1	300	0.5578 +/- 0.0321	0.9824
4	20	1	300	0.5562 +/- 0.0354	0.9994
5	10	2	300	0.5512 +/- 0.0284	0.8241

Parameter Selection Rationale:

- **Tree Depth (`max_depth = 15`):** A maximum depth of 15 captures non-linear chemical thresholds without letting individual trees memorize noise. Restricting depth to 10 resulted in underfitting ($F1 = 0.5512$), while depth 20 raised training F1 (0.9615) but slightly lowered validation F1 (0.5584), showing early signs of overfitting.
- **Leaf Regularization (`min_samples_leaf = 2`):** Requiring at least 2 samples per leaf prevents terminal nodes from fitting to single noisy samples, stabilizing validation performance.
- **Ensemble Scale (`n_estimators = 300`):** An ensemble of 300 trees provides sufficient variance reduction, stabilizing probability estimates and feature importance rankings.

Final Model Choice

Based on cross-validation results, statistical hypothesis testing, and hyperparameter tuning, we selected the **Random Forest Classifier** (`n_estimators = 300`, `max_depth = 15`, `min_samples_split = 4`, `min_samples_leaf = 2`, `class_weight = 'balanced'`) as our final model. It provides:

1. Statistically significant outperformance over the linear baseline ($p = 0.0081$ uncorrected, $p = 0.0448$ Nadeau & Bengio corrected).
2. The highest validation F1-score (0.5599), ROC-AUC (0.8492), and PR-AUC (0.5747).
3. Inherent invariance to monotonic transformations and outliers.
4. Reliable rank-order discrimination suitable for decision threshold tuning.

Section 5 presents the out-of-sample evaluation and diagnostic analysis of this model on the held-out test set.

5. Model Analysis, Diagnostics, and Operational Evaluation

We evaluated our final Random Forest model on the 1,064 unseen wines in `data/test.csv`. This section examines out-of-sample test performance, error trade-offs, diagnostic curves, which chemical features most strongly drove predictions, and how the model performed on red versus white wines.

Test Set Performance

Table 8 summarizes the final performance metrics evaluated on the 1,064 held-out test observations.

Table 8. Final Random Forest Evaluation on Held-Out Test Dataset (N = 1,064)

Performance Metric	Test Score	Operational Interpretation
Accuracy	83.08% (0.8308)	Correctly classifies 884 of 1,064 total test wines
Precision	54.10% (0.5410)	145 of 268 wines flagged as high quality are true premium selections
Recall (Sensitivity)	71.78% (0.7178)	Successfully identifies 145 of the 202 true high-quality test wines
Specificity	85.73% (0.8573)	Correctly identifies 739 of 862 ordinary wines as ordinary
F1-Score	0.6170	Robust harmonic balance between precision and sensitivity
ROC-AUC	0.8738	Strong global discriminability across all operating thresholds
PR-AUC (Average Precision)	0.6230	Substantial minority-class enrichment (+43.32% over 18.98% baseline)
Brier Score Loss	0.1225	Low quadratic probability error reflecting solid calibration

Table 9 presents the detailed classification report for both target classes.


Table 9. Test Classification Report by Class

Class Label	Support (N)	Precision	Recall	F1-Score
Ordinary Wine (<code>is_good = 0</code>)	862 (81.02%)	0.9284	0.8573	0.8914
High-Quality Wine (<code>is_good = 1</code>)	202 (18.98%)	0.5410	0.7178	0.6170
Macro Average	1,064	0.7347	0.7876	0.7542
Weighted Average	1,064	0.8549	0.8308	0.8393

The test set results demonstrate strong generalization. Test F1-score (0.6170) and test ROC-AUC (0.8738) exceed the 10-fold cross-validation estimates (0.5599 and 0.8492, respectively), confirming that the model did not suffer from overfitting during cross-validation.

Confusion Matrix and Decision Trade-offs

Figure 8 displays the raw count and normalized confusion matrices on the test dataset.



figure_02_confusion_matrix.png

Out of 1,064 test evaluations:

- **True Negatives (TN):** 739 samples (69.45% of total, 85.73% of ordinary wines).
- **True Positives (TP):** 145 samples (13.63% of total, 71.78% of high-quality wines).
- **False Positives (FP: Type I Error):** 123 samples (11.56% of total, 14.27% of ordinary wines).
- **False Negatives (FN: Type II Error):** 57 samples (5.36% of total, 28.22% of high-quality wines).

Practical Implications of Errors:

1. Type I Errors (False Positives: 123 wines):


- *Operational Consequence:* An ordinary wine is labeled as high quality and could be routed to a premium bottling line.
- *Risk:* Subpar wine on store shelves under a reserve label damages consumer trust and brand reputation.
- *Mitigation:* The model can serve as a pre-screening filter, with flagged wines confirmed by human sensory evaluation before bottling.

2. Type II Errors (False Negatives: 57 wines):

- *Operational Consequence:* A high-quality wine is labeled ordinary and sold in bulk or standard distribution.
- *Risk:* Forfeited margin. Premium wine sold as table wine loses potential revenue.
- *Mitigation:* The balanced class weighting keeps false negatives low (only 5.36% of total samples), catching over 71.7% of all premium lots.

ROC, Precision-Recall, and Calibration Curves

Figure 9 illustrates the Receiver Operating Characteristic (ROC) curve alongside the Precision-Recall (PR) curve for the held-out test dataset.




figure_03_roc_and_pr_curves.png

Key Diagnostic Curve Observations:

- **ROC Trajectory (AUC = 0.8738):** The ROC curve ascends steeply in the low false positive rate region (0.00 to 0.20), reaching over 70% recall while keeping the false positive rate below 15%. This demonstrates strong discriminative power across wide operating conditions.
- **Precision-Recall Dynamics (PR-AUC = 0.6230):** In imbalanced domains, the PR curve provides an informative assessment of minority class precision. The horizontal reference line at 0.1898 represents a random guess baseline. The Random Forest model maintains a substantial elevation above this baseline across the entire recall spectrum.
- **Adjustable Decision Thresholds:** The default 0.50 probability threshold yields 54.10% precision at 71.78% recall (F1 = 0.6170). Depending on commercial strategy, the operating threshold can be adjusted along the precision-recall trade-off continuum:
 - *Brand Protection Policy (High Precision):* Elevating the classification threshold to 0.65 increases precision to 62.86% (at 43.56% recall), while setting the threshold to 0.70 yields 68.93% precision (at 35.15% recall) and 0.75 yields 73.68% precision (at 27.72% recall), ensuring that only high-confidence lots receive premium bottling.
 - *Volume Harvesting Policy (High Recall):* Lowering the threshold to 0.40 captures 80.69% of true high-quality wines (at 45.79% precision), while setting the threshold to 0.35 captures 85.15% of high-quality wines (at 41.95% precision) for secondary sommelier auditing.

Figure 10 displays the probability calibration curve (reliability diagram).



figure_06_calibration_curve.png

The calibration curve illustrates a characteristic upward probability shift: across the lower and middle probability intervals, predicted probabilities exceed empirical event frequencies (for example, wines assigned a model probability of ~0.45 exhibit an actual high-quality rate of 20.45%). This probability distortion is a direct mathematical consequence of cost-sensitive learning (`class_weight='balanced'`), which artificially scales minority class misclassification loss by a factor of 2.64 during tree splitting. While this reweighting significantly improves minority class discrimination and decision-boundary separability (achieving a test ROC-AUC of 0.8738 and F1-score of 0.6170), raw model probabilities should be interpreted as relative risk indices rather than uncalibrated Bayesian posterior probabilities unless post-hoc calibration (such as Platt scaling or isotonic regression) is applied. Despite this systematic shift, the overall test Brier score loss is 0.1225 (substantially outperforming a naive prevalence baseline of 0.1538).

Feature Importance

To understand the chemical drivers governing wine quality classification, we analyzed feature importance using two distinct analytical methods:


- Gini Impurity Importance (Mean Decrease in Impurity, MDI):** Total impurity reduction across all 300 bagged trees during training.
- Permutation Feature Importance:** Model-agnostic drop in test F1-score when each feature is randomly shuffled across 20 iterations on `data/test.csv`.

Table 10 compares the rankings from both methodologies.

Table 10. Physicochemical Feature Importance Comparison

Feature Name	Permutation Mean (F1 Drop)	Permutation Std	Gini Impurity (MDI)	Physical/Chemical Mechanism
Alcohol	0.2461	0.0268	0.2251	Ethanol content: body, perceived sweetness, aromatic volatility
Chlorides	0.0796	0.0167	0.0869	Salt content: mineral mouthfeel; excess imparts saltiness and bitterness
Density	0.0768	0.0151	0.1235	Inversely related to alcohol and sugar: overall mouthfeel and body
Total Sulfur Dioxide	0.0549	0.0177	0.0799	Preservative: antioxidant/antimicrobial; excess dulls fruit flavors
Volatile Acidity	0.0437	0.0146	0.0858	Acetic acid: vinegar fault; lower levels required for clean sensory profile
Sulphates	0.0432	0.0093	0.0696	Potassium sulphate: antimicrobial synergy, fruit freshness enhancer
Free Sulfur Dioxide	0.0377	0.0113	0.0652	Active unbound preservative protecting wine from oxidation
Citric Acid	0.0367	0.0137	0.0738	Minor organic acid: provides crispness and fresh citrus notes
pH	0.0323	0.0085	0.0613	Acidity buffer: microbial stability and perception of tartness
Residual Sugar	0.0256	0.0082	0.0678	Unfermented sugar: balance against high acidity in Vinho Verde
Fixed Acidity	0.0250	0.0072	0.0576	Tartaric and malic acids: foundational crisp structural acidity
is_red	0.0044	0.0040	0.0035	Variety indicator: secondary after direct chemical measurements

Figure 11 illustrates the comparative importance profiles.



figure_04_feature_importance.png

Key Chemical Insights:

- 1. Alcohol is the Dominant Quality Predictor:** Shuffling alcohol on the test set drops the model's F1-score by 0.2461 (from 0.6170 down to ~0.3709). Gini importance confirms this finding (0.2251). In cool maritime climates like Vinho Verde, higher natural alcohol indicates fully ripe grapes with concentrated flavors, balanced acids, and proper phenolic maturity.
- 2. Chlorides and Density Form Key Structural Markers:** Chlorides (salt content) and density represent the second and third most impactful features in permutation testing. Excessive chlorides impart an unpleasant saline astringency that judges penalize heavily. Density integrates alcohol and sugar levels, serving as a composite indicator of extraction and body.
- 3. Acid Balance and Fermentation Cleanliness:** Volatile acidity (acetic acid) is a primary marker of bacterial spoilage. High-quality wines consistently maintain low volatile acidity. Sulphates and sulfur dioxide provide microbial protection, ensuring fruit notes remain clean and stable.
- 4. Wine Type (`is_red`) Has Minor Direct Impact:** The binary type indicator shows low permutation importance (0.0044). Once specific chemical concentrations (acidity, density, and sulfur levels) are known, the explicit red vs. white label provides little additional predictive information.


Red vs. White Wine Performance

To evaluate whether the model functions equally well across wine varieties, we disaggregated test set performance by wine type.

Table 11. Subgroup Disaggregation on Test Dataset

Wine Variety	Samples (N)	Good Wines (N)	Prevalence	Accuracy	Precision	Recall	F1-Score	ROC-AUC	PR-AUC
Red Wine	275	37	13.45%	90.18%	60.42%	78.38%	0.6824	0.9307	0.7255
White Wine	789	165	20.91%	80.61%	52.73%	70.30%	0.6026	0.8522	0.6044

Figure 12 shows the disaggregated ROC and PR curves for red and white wines.



figure_05_subgroup_analysis.png

Subgroup Findings:

- **Red Wine Results:** The model achieves higher performance on red wines: 90.18% accuracy, 78.38% recall, 0.6824 F1-score, and an ROC-AUC of 0.9307 (with PR-AUC of 0.7255 against a 13.45% baseline). In red wines, high quality is strongly demarcated by low volatile acidity (avoiding vinegar spoilage) and high sulphates.
- **White Wine Results:** White wines exhibit an ROC-AUC of 0.8522 and an F1-score of 0.6026. White wines are more heterogeneous in residual sugar and acidity balance (ranging from dry to off-dry), creating more overlap between ordinary and high-quality profiles.
- **Takeaway:** The model provides strong classification utility across both styles, with especially high discrimination on red wines.

Model Validity and Overfitting Check

To verify algorithmic stability and confirm that the model did not memorize training artifacts, we compared performance across training, 10-fold cross-validation, and held-out testing partitions.

Table 12. Performance Comparison across Training, Cross-Validation, and Test Sets

Metric	Training Set (Fit, N = 4,256)	10-Fold CV Validation (Mean +/- Std)	Held-Out Test Set (Unseen, N = 1,064)
Accuracy	97.46%	80.50% +/- 1.70%	83.08%
F1-Score	0.9372	0.5599 +/- 0.0338	0.6170
ROC-AUC	0.9991	0.8492 +/- 0.0155	0.8738
Precision	0.8924	0.4908 +/- 0.0352	0.5410
Recall	0.9864	0.6543 +/- 0.0501	0.7178

Findings:

1. **Regularization Effectiveness:** While tree ensembles naturally exhibit high training fit due to bootstrap sample isolation, the validation metrics across 10 folds remained stable with low variance ($\sigma_{\text{ROC-AUC}} = 0.0155$, $\sigma_{\text{F1}} = 0.0338$).
2. **Generalization Confirmation:** The test set scores (F1 = 0.6170, ROC-AUC = 0.8738) closely mirror and slightly exceed the cross-validation estimates. This confirms that the hyperparameter constraints (`max_depth = 15`, `min_samples_leaf = 2`) prevented the ensemble from overfitting.

Limitations

While the Random Forest classifier demonstrates strong empirical performance, practical deployment in commercial viticulture must account for several structural limitations:

1. **Regional viticultural specificity:** All samples originate from the Vinho Verde region in northwest Portugal. Vinho Verde is characterized by a cool, humid Atlantic climate producing wines with high natural acidity, lower alcohol levels, and distinctive carbonation. The physicochemical boundaries learned by this model may not generalize to warmer viticultural regions (such as Napa Valley, Rioja, or Barossa Valley) where baseline alcohol, sugar, and tannin profiles differ markedly.
2. **Subjectivity in sensory quality ratings:** The ground-truth quality score represents the median rating from at least three sensory experts on a 0 to 10 scale. Expert wine tasting involves subjective palate differences, fatigue during lengthy tasting flights, and ambient temperature sensitivity. This sensory variance introduces irreducible label noise near the decision boundary (qualities 6 vs. 7).

3. **Absence of commercial and viticultural variables:** The dataset contains purely post-fermentation laboratory chemistry. It omits critical factors that heavily drive commercial market pricing, including:

- Specific grape cultivar (e.g., Alvarinho, Loureiro, Vinhão).
- Vintage year and meteorological growing season conditions.
- Oak barrel maturation duration and toast profiles.
- Phenolic and anthocyanin compounds (tannins, color pigments).
- Brand equity, packaging, and retail distribution channels.

Business Recommendations

The model evaluation demonstrates that laboratory physicochemical measurements can accurately identify high-quality wines (ROC-AUC = 0.8738, F1 = 0.6170).

For winemakers and commercial distributors, we recommend:

1. **Automated Screening:** Deploy the model as an automated batch screening tool before formal sensory evaluation, triaging production volume while flagging high-potential lots.
2. **Physicochemical Monitoring:** Actively monitor and manage alcohol concentration, volatile acidity levels, and chloride balance during fermentation and blending.
3. **Threshold Calibration:** Adjust operational decision thresholds based on business priorities, using conservative thresholds (0.70 to 0.75) for premium reserve bottling and sensitive thresholds (0.35 to 0.40) for comprehensive quality audits.

Discussion

This discussion summarizes the empirical findings, enological mechanics, and operational implications derived from the complete modeling lifecycle (Sections 4 and 5).

Model Performance and Architecture

Across 10-fold stratified cross-validation and independent out-of-sample testing on 1,064 unseen wines, the bagging Random Forest classifier demonstrated consistent superiority over baseline and comparative architectures:

- **Baseline Outperformance:** Both the standard paired Student's t-test ($t = 3.3840$, $p = 0.0081$) and the Nadeau and Bengio (2003) corrected resampled t-test ($t = 2.3290$, $p = 0.0448$) confirmed a statistically significant improvement over Logistic Regression at $\alpha = 0.05$. Linear models exhibited high recall (78.69%) but poor precision (40.21%), generating an unacceptable rate of false positives.
- **Ensemble Robustness:** Constrained tree depth (`max_depth = 15`) and leaf regularization (`min_samples_leaf = 2`) effectively mitigated overfitting, evidenced by test metrics (83.08% accuracy, 0.8738 ROC-AUC, 0.6170 F1-score) closely matching 10-fold cross-validation estimates (80.50% accuracy, 0.8492 ROC-AUC, 0.5599 F1-score).
- **Literature Benchmark Context:** Cortez et al. (2009) originally modeled wine quality as a continuous regression task on the 0 to 10 scale, evaluating support vector machines via Mean Absolute Deviation (MAD) and Regression Error Characteristic (REC) curves. Because this investigation reframes the task as an operational binary classification problem (quality score 7 or higher vs. 6 or lower), direct numerical comparison to the original regression error is not applicable; however, our Random Forest classifier establishes a high-performance enological classification benchmark, achieving a test ROC-AUC of 0.8738 and a red wine subgroup ROC-AUC of 0.9307.

Key Chemical Drivers

Feature importance evaluations across both in-sample Gini impurity and out-of-sample permutation testing reveal three primary enological mechanisms:

1. **Ethanol as the Primary Maturity Marker:** Alcohol content emerged as the single most critical predictor (causing a \$0.2461\$ drop in test F1-score under permutation). In the cool Atlantic climate of the Vinho Verde region, achieving higher natural alcohol ($> 11.0\%$ ABV) signifies prolonged hang-time, optimal grape maturity, and concentrated flavor precursors without excess harsh acidity.
2. **Structural Balance via Density and Chlorides:** Density (\$+0.0768\$ F1 drop) and chlorides (\$+0.0796\$ F1 drop) function as secondary structural constraints. Elevated chlorides impart an unpleasant saline or soapy character that sensory panels penalize severely, while density captures the ratio of residual extract to alcohol.
3. **Microbial Cleanliness via Volatile Acidity and Sulphates:** Volatile acidity (acetic acid) acts as an acute negative indicator. High-quality wines consistently maintain low volatile acidity, supported by adequate free and total sulfur dioxide preservation.

Red vs. White Wine Performance

Disaggregating model performance revealed distinct predictive dynamics:

- **Red Wines (\$90.18\%\$ Accuracy, \$0.9307\$ ROC-AUC):** Red wines exhibited superior separability because defect markers (excess volatile acidity) and preservation markers (sulphates) provide sharp boundaries.
- **White Wines (\$80.61\%\$ Accuracy, \$0.8522\$ ROC-AUC):** White wines present broader variation in residual sugar and acidity balance (ranging from dry to semi-sweet styles), creating greater overlap between quality classes.
- Explicit wine type (`is_red`) demonstrated low permutation importance (\$+0.0044\$), confirming that once specific chemical concentrations are known, the categorical color designation offers minimal additional predictive utility.

Performance Limits

Extensive diagnostic testing demonstrated that the model operates near the empirical performance ceiling of this dataset:

- **Threshold Sensitivity:** Scanning decision thresholds from \$0.30\$ to \$0.74\$ demonstrated that the default \$0.50\$ cutoff (\$F1 = 0.6170\$) is within \$0.0025\$ of the empirical maximum (\$F1 = 0.6195\$ at threshold \$0.52\$).
- **Interaction Ratios:** Testing domain-engineered chemical interaction ratios (such as free-to-total SO2 ratio, total acidity, and alcohol-to-density ratio) did not improve cross-validation performance (\$F1 = 0.5587\$ vs. \$0.5599\$ baseline), confirming that tree-based algorithms natively partition non-linear interactions without redundant synthetic features.
- **Sensory Noise:** The remaining test classification error (\$16.92\%\$) likely reflects a combination of subtle non-linear chemical dynamics and inherent sensory noise from human expert evaluation (which relies on subjective palate ratings across multi-taster panels), rather than simple algorithmic underfitting.

Practical Recommendations

For commercial wine producers, distributors, and quality assurance laboratories, we recommend a three-tiered operational deployment:

1. **Tier 1 (Automated Pre-Screening):** Run laboratory chemical measurements through the model immediately post-fermentation to triage production batches into standard table wine vs. potential reserve quality.
2. **Tier 2 (Risk-Adjusted Decision Thresholds):** Calibrate the operating threshold to match commercial objectives: use a high-precision threshold (≥ 0.70) for premium reserve bottling to protect brand reputation, and a high-recall threshold (≤ 0.40) for bulk screening to capture high-potential lots.
3. **Tier 3 (Targeted Expert Sensory Flights):** Reserve high-cost human sommelier panels specifically for lots flagged by the model as borderline or high-probability premium. Because only \$25.19\%\$ of test batches are flagged for premium evaluation under the default threshold (\$268\$ of \$1,064\$ wines), this triage protocol could theoretically

reduce routine tasting volume by approximately 75%, allowing winemakers to focus palate evaluations where they add the greatest economic value.

6. References (APA 7th Edition)

- Cortez, P., Cerdeira, A., Almeida, F., Matos, T., & Reis, J. (2009). Modeling wine preferences by data mining from physicochemical properties. *Decision Support Systems*, 47(4), 547-553. <https://doi.org/10.1016/j.dss.2009.05.016>
- Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5-32. <https://doi.org/10.1023/A:1010933404324>
- Nadeau, C., & Bengio, Y. (2003). Inference for the generalization error. *Machine Learning*, 52(3), 239-281. <https://doi.org/10.1023/A:1024068626366>
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.
- UCI Machine Learning Repository. (2009). *Wine quality dataset* [Data set]. University of California, Irvine. <https://archive.ics.uci.edu/dataset/186/wine+quality>